

VIGIL

Evidence before conclusions.

An accountable incident-response layer for Indian bank SOC teams
—built on Elastic, designed to stop at human judgment.

TWO SIX-HOUR CLOCKS

DAKSH

from detection

CERT-In

from noticing

Draft · review · evidence

Two clocks. Different anchors.

The prototype prepares evidence; an authorized human remains responsible for reporting.

01 RBI DAKSH

Within six hours of detection

Applies to banks covered by the 2026 commercial-bank Directions. The same provision also requires proactive CERT-In notification.

02 CERT-In

Within six hours of noticing

For incident types listed in Annexure I—including unauthorized access and incidents affecting digital-payment systems.

VIGIL does not auto-file. It builds a reviewable evidence package and preserves the analyst decision.

What the rehearsal actually proved

One fixed synthetic scenario, one saved dashboard, one inspectable path.

136

synthetic ECS-aligned events

3

final behaviour alerts

1

potential discovery

₹11.6L

UPI fixture value · 5 accounts

12 CORRELATED EVIDENCE EVENTS

VIGIL · SYNTHETIC PAYMENT COMPROMISE EVIDENCE

Evidence before conclusions

Sanitized public view · live environment capture retained privately

12

correlated events

₹11,60,000

suspicious UPI fixture


5

affected accounts

Synthetic evidence timeline


Attack Discovery is a lead—not a verdict

VIGIL keeps the speed of AI and makes certainty earn its way into the case.



One workflow, accountable by design

Native Elastic primitives do the platform work; VIGIL adds bank-specific controls around them.



Fail closed

Missing evidence, unsupported certainty, or absent human input stops the path.
Destructive and external actions stay disabled in the rehearsal.

The human gate is real

A controlled full-chain run paused, accepted HOLD, and completed without an external action.



Evidence	Execution and Case IDs retained in the private rehearsal log	70.9s CONTROLLED RUN
Public deck	Environment identifiers removed before repository publication	21,852 LLM tokens

Synthetic, but structurally honest


We preserve the interfaces the product needs without pretending to possess bank production data.

LAYER	WHAT WE MODEL	HONEST BOUNDARY
 Security telemetry	ECS-aligned events: auth, process/network context, alerts	Not a claim of production volume or a fraud model
 Bank context	Explicit lookup fields for account, tier, UPI batch, amount	Not NPCI wire format; validate per bank integration
 Regulator draft	Fields from CERT-In’s public Incident Reporting Form + DAKSH pack	Review-only; not an official filing or portal submission

STIX/TAXII: optional threat-intelligence exchange—not the public CERT-In incident-report format.

Evidence can be challenged

The prototype proves tamper detection locally; it does not claim immutable storage.



Prototype boundary: SHA-256 hash chain + standalone verifier. S3 Object Lock and KMS signing are future controls until demonstrated.

From AI lead to accountable decision

Three beats. Ninety seconds. Every claim visible on screen.

01

GROUND

Dashboard + ES|QL
12 events · ₹11.6L · 5 accounts

02

CHALLENGE

Flag unsupported certainty
Keep source alerts attached

03

DECIDE ♦ VERIFY

Human HOLD
Genuine PASS · tampered FAIL

EVENT-DAY TRUTH

Rebuild in the official repo, disclose prior rehearsal, and prove an AWS-hosted LLM path before claiming compliance.

The rehearsal removes platform risk; it does not replace the judged build.